



ParaPerformance: An Advanced Refactoring Tool for Parallelising C++ Programs

Chris Brown

Craig Manson, Kenneth MacKenzie, Vladimir Janjic, Kevin Hammond

University of St Andrews, Scotland

@chrismarkbrown

@rephrase_eu

cmb21@st-andrews.ac.uk

C++ Meetup Edinburgh – November 2015

RE²PHRASE
RE²SHRASE

ParaPerformance



Scottish Enterprise



RePhrase Project: Refactoring Parallel Heterogeneous Software – a Software Engineering Approach (ICT-644235), 2015-2018, €3.6M budget

8 Partners, 6 European countries
UK, Spain, Italy, Austria, Hungary, Israel



UNIVERSITÀ DI PISA

UNIVERSITÀ
DEGLI STUDI
DI TORINO
ALMA UNIVERSITAS
TAURINENSIS

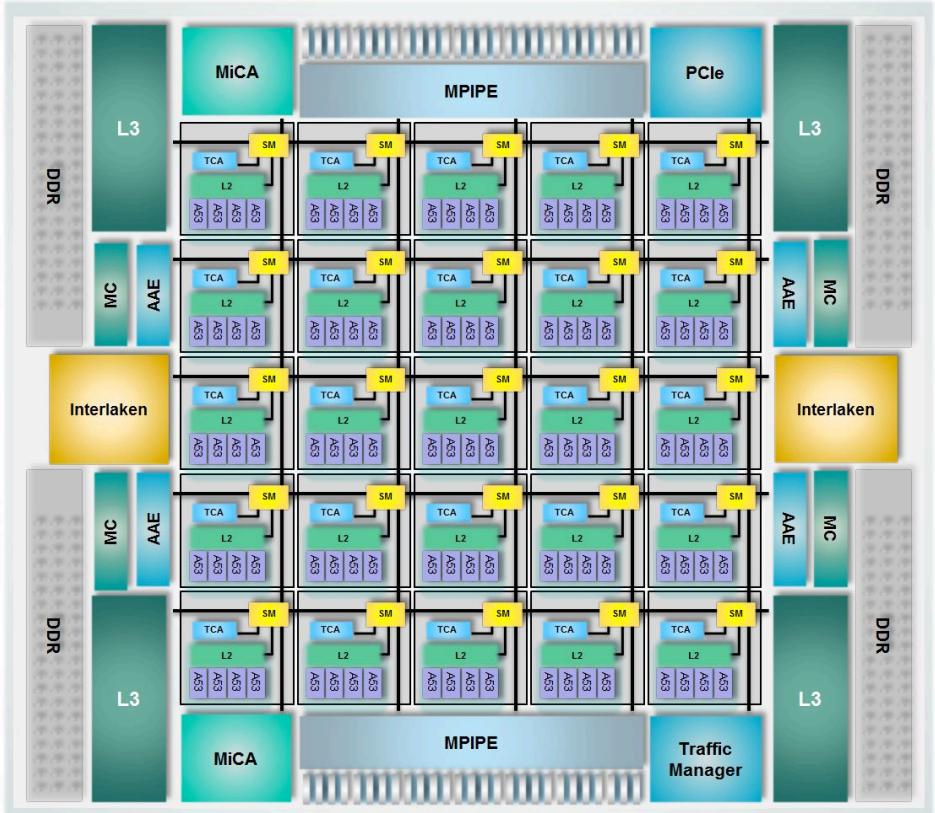


ParaFormance

The ParaFormance Project

- Seed Funding from Scottish Enterprise
 - 1-3 years pre-commercialisation funding
 - Develop work from EU/UK publicly-funded projects
- Start Date
 - 1st June 2015
- Led by: Kevin Hammond and Chris Brown, University of St Andrews

The Dawn of a New Age



- EZCHIP – TILE-MX100
- 100 64-bit AMD x86 Cores
- 3-level cache with > 40 Mbytes on-chip cache
- Multitude of network accelerators
- Over 200Gbps integrated I/O including Ethernet,
- DDR supports up to 1 TB RAM

It's not just about large systems

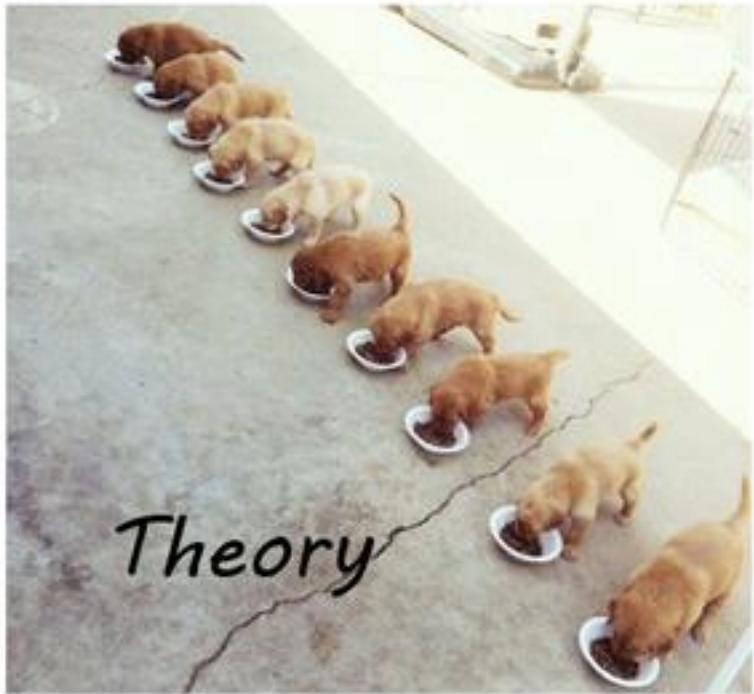
- Even mobile phones are multicore
 - Samsung Exynos 5 Octa has 8 cores, 4 of which are “dark”
- Performance/energy tradeoffs mean systems will be increasingly parallel
- If we don't solve the multicore challenge, then no other advances will matter!



ALL Future
Programming will be
Parallel!

Programming Multicore Systems...

Multithreaded programming



Theory



Actual

REPHRASE
BE SHREASE

ParaFormance

Thinking in Parallel

- Fundamentally, programmers must learn to “*think parallel*”
 - this requires new *high-level* programming constructs
 - you cannot program effectively while worrying about deadlocks etc.
 - **they must be eliminated from the design!**
 - you cannot program effectively while fiddling with communication etc.
 - **this needs to be packaged/abstracted!**
 - you cannot program effectively without performance information
 - **this needs to be included!**
- We use two key technologies:
 - Refactoring (changing the source code structure)
 - Parallel Patterns (high-level functions of parallel algorithms)

RE  **PHRASE**
BE  **SHRASE**

ParaPerformance

Parallel Patterns

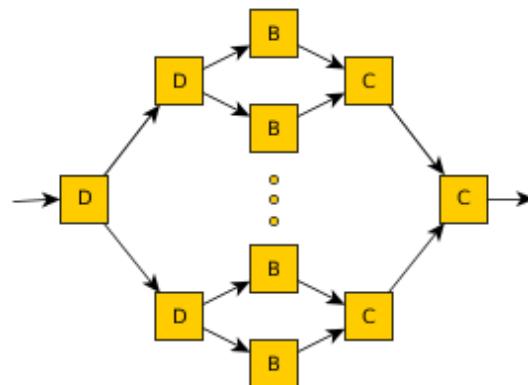


The diagram illustrates the Farm architecture. A central yellow square node labeled "W" is connected to two yellow circular nodes labeled "W". The left "W" node has three incoming arrows from the bottom-left, and two outgoing arrows pointing to the right "W" node. The right "W" node also has two incoming arrows from the left "W" node and one outgoing arrow pointing to the right. Below the left "W" node is the text "Schedule tasks", and below the right "W" node is the text "Gather results".

A horizontal sequence of three yellow rectangular boxes connected by arrows. The first box contains 'S1', the second 'S2', and the third 'Sn'. An arrow points from the left towards 'S1', and another arrow points from 'Sn' towards the right.

The diagram illustrates the Reduce phase of MapReduce. It shows multiple mappers (represented by yellow squares labeled 'f') receiving input from multiple sources and producing output. The outputs are then aggregated by a reducer (another yellow square labeled 'f').

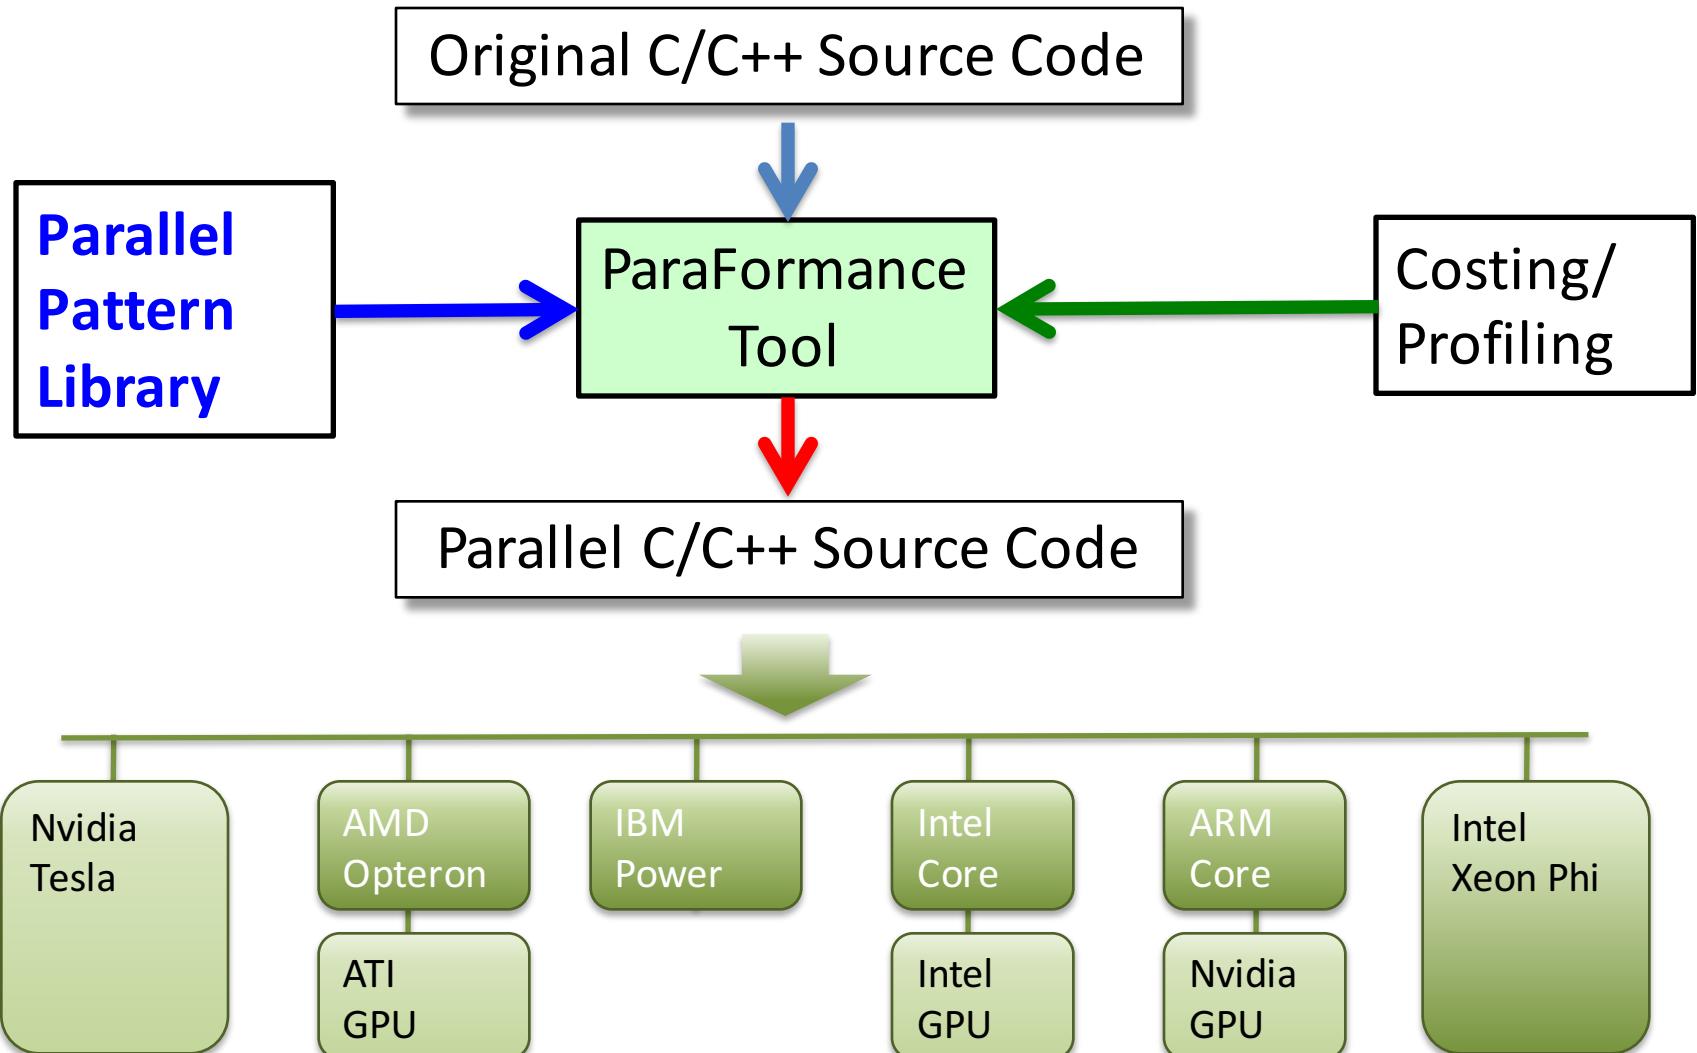
Divide&Conquer



Our Approach

- Start bottom-up
 - Identify components (side-effect free functions that correspond to parallel computations)
 - *using semi-automated refactoring*
- Think about the **PATTERN** of parallelism
 - e.g. map(reduce), task farm, parallel search, parallel completion, ...
- Structure the components into a parallel program
 - *turn the patterns into concrete (skeleton) code*
 - Take performance, energy etc. into account (multi-objective optimisation)
 - *also using refactoring*
- Restructure if necessary!
 - *also using refactoring*

General Technique



Sequential Refactoring

Semi-automated (user-driven)

1. Renaming
2. Inlining
3. Changing scope
4. Adding arguments
5. Generalising Definitions
6. Type Changes



Examples include refactoring Linux kernels using Coccinelle, refactoring Java/C++ in Eclipse, etc.

Refactoring = Condition + Transformation

Transformation

- Ensure change at all points needed.
- Ensure change at only those points needed.

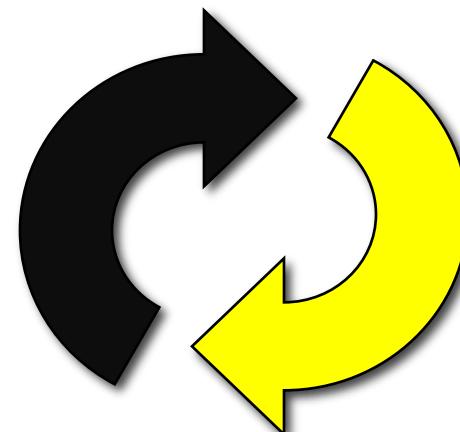
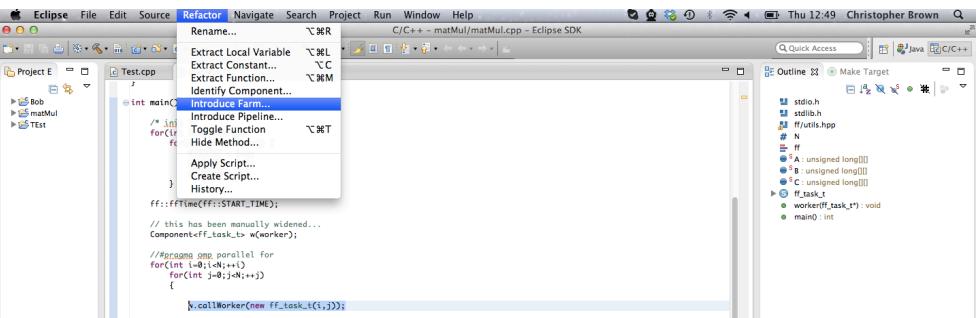
Condition

- Is the refactoring applicable?
- Will it preserve the “semantics” of the program?
- The module? The File?

Both pre- and post- conditions

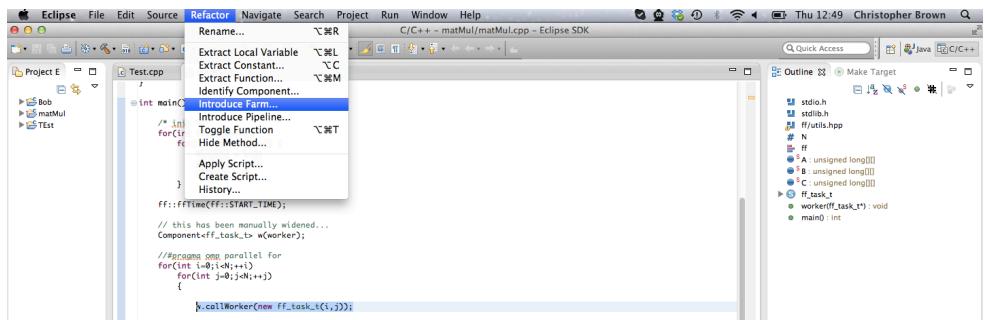
Refactoring can help parallel thinking!

- can be used to introduce parallelism, and help choose the right abstraction
- parallel programs can be refactored for new parallel architectures
- can check the conditions for applying parallel skeletons
- performance information can be integrated
- Programmer ‘in the loop’



The ParaFormance Tool Prototype

- Integrated into Eclipse (CDT)
- Supports full C++(11) standard
- Layout and comment preserving
- Undoable
- Preview feature



REPHRASE
RESHARAE

ParaFormance

Our Refactorings

- Refactorings to Introduce:
 - Farm/Map, Parallel-For and Pipeline patterns
 - FastFlow
 - Components
 - Farm
 - Pipeline
 - TBB
 - Lambda
 - Function Class
 - Parallel For/Pipeline
 - OpenMP
 - Parallel-For

The ParaFormance Toolkit

1. Prediction/Estimation
 - Accurate and advance performance modelling
2. Discovery
 - Automatic discovery of (instances of) parallel patterns
3. Insertion
 - Automatically insert the parallel “Business Logic”
4. Elimination
 - Remove existing/legacy parallelism
5. Validity
 - Fundamental condition checks
6. Profile
7. Shaping



Image Convolution

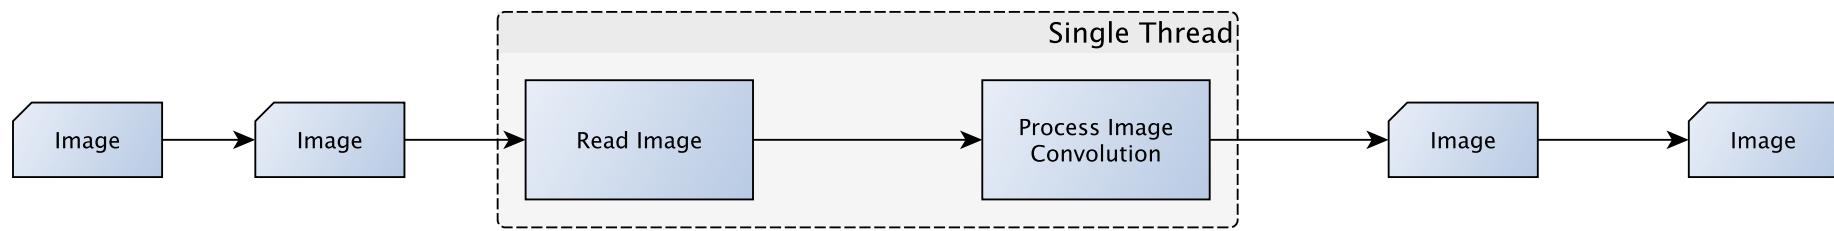
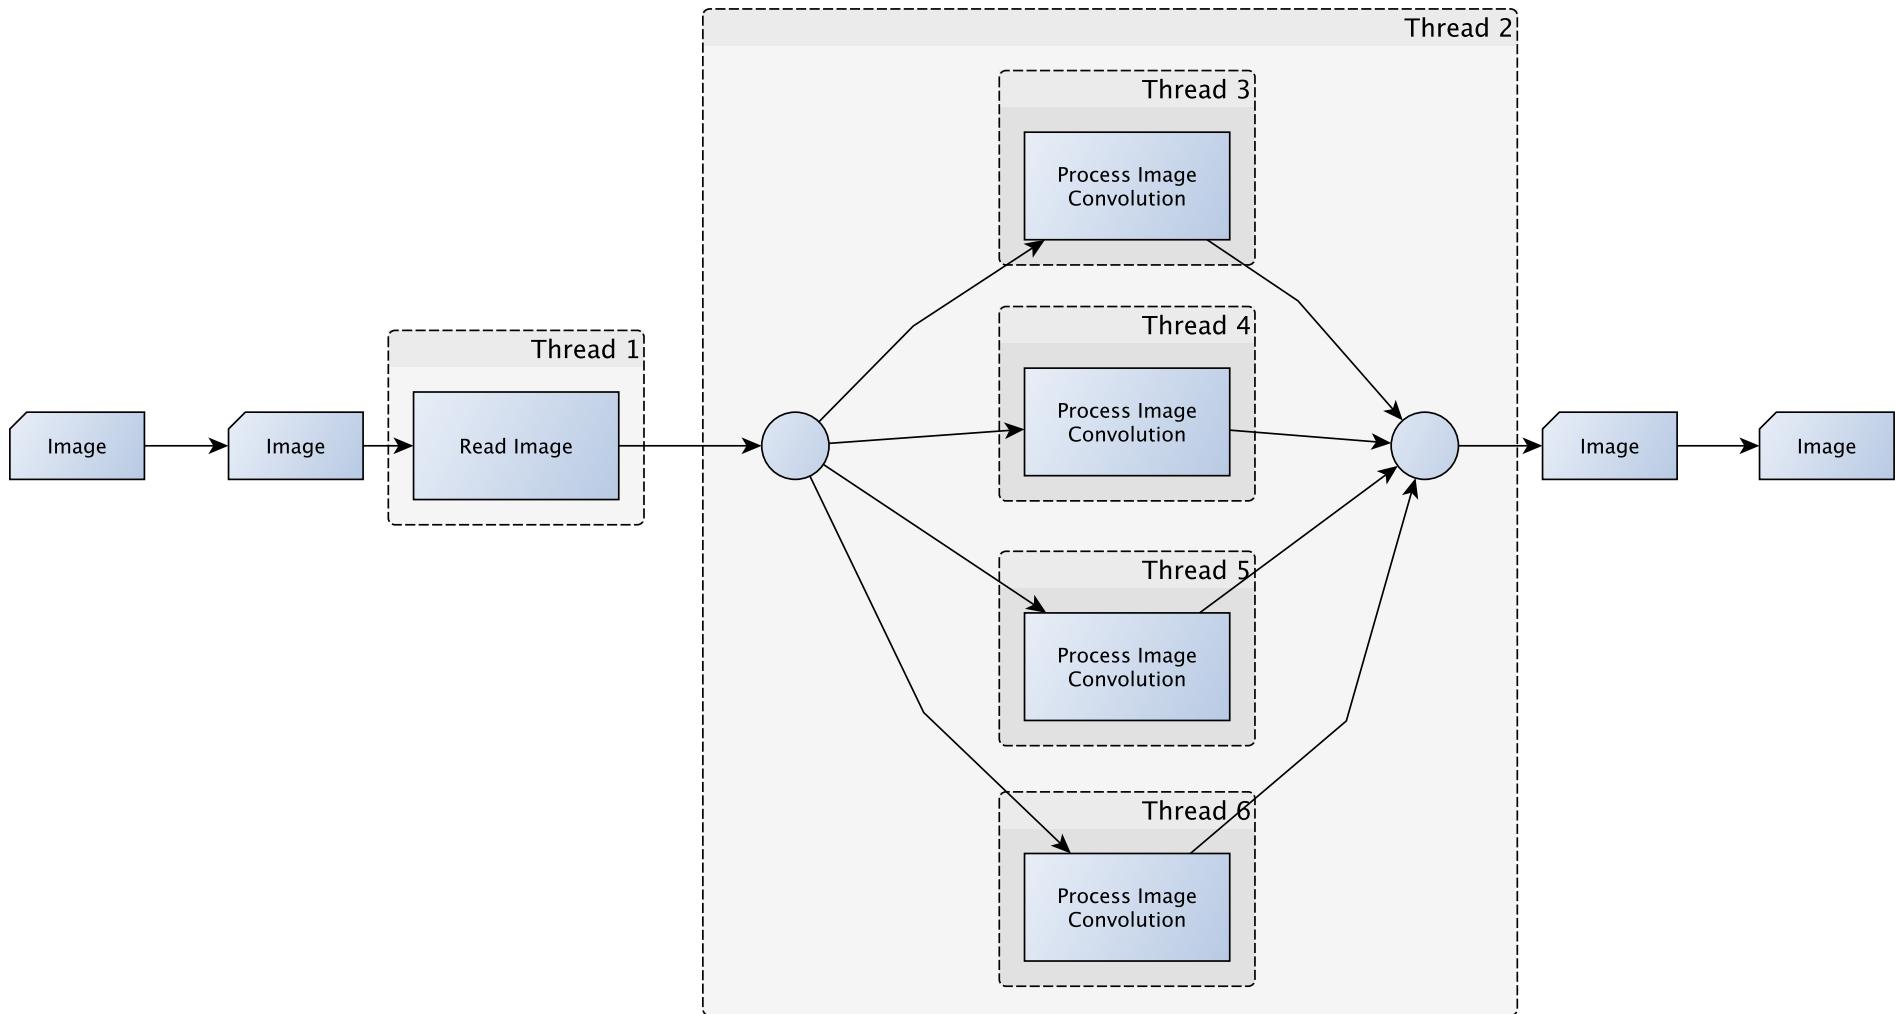


Image Convolution, Refactored





University
of
St Andrews

C++ Refactoring Demo

REPHRASE
REPHRASE

ParaPerformance

Image Convolution

Speedups for Image Convolution on *xookik*

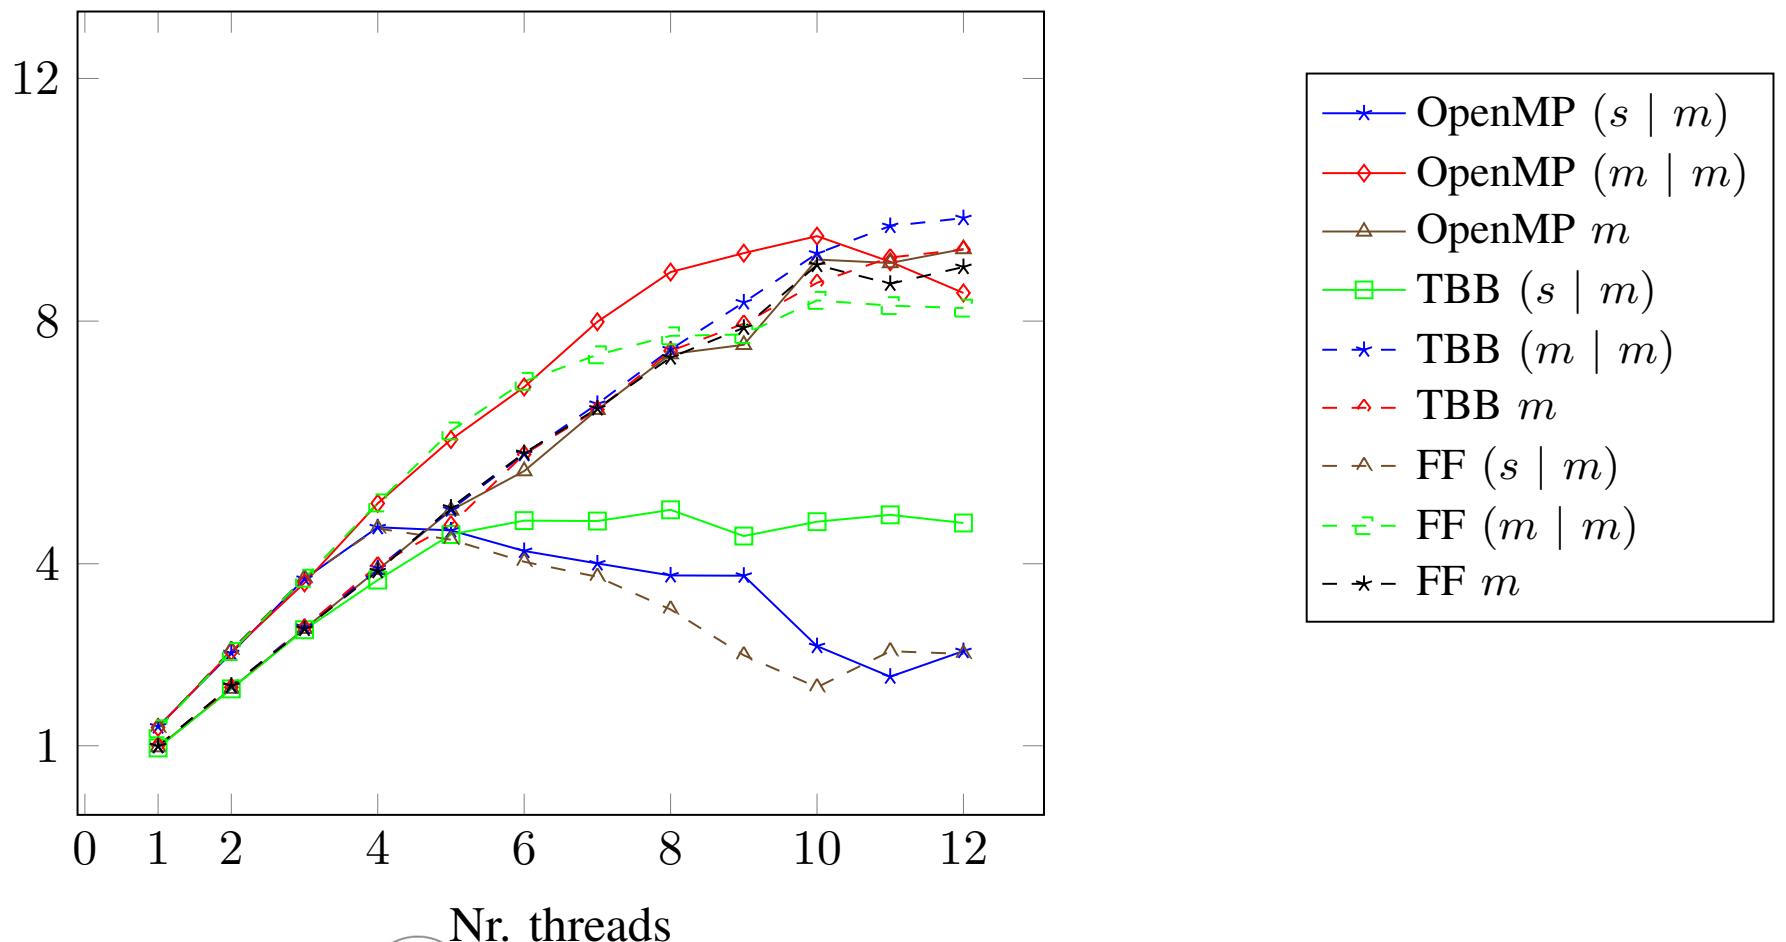


Image Convolution

Speedups for Image Convolution on *titanic*

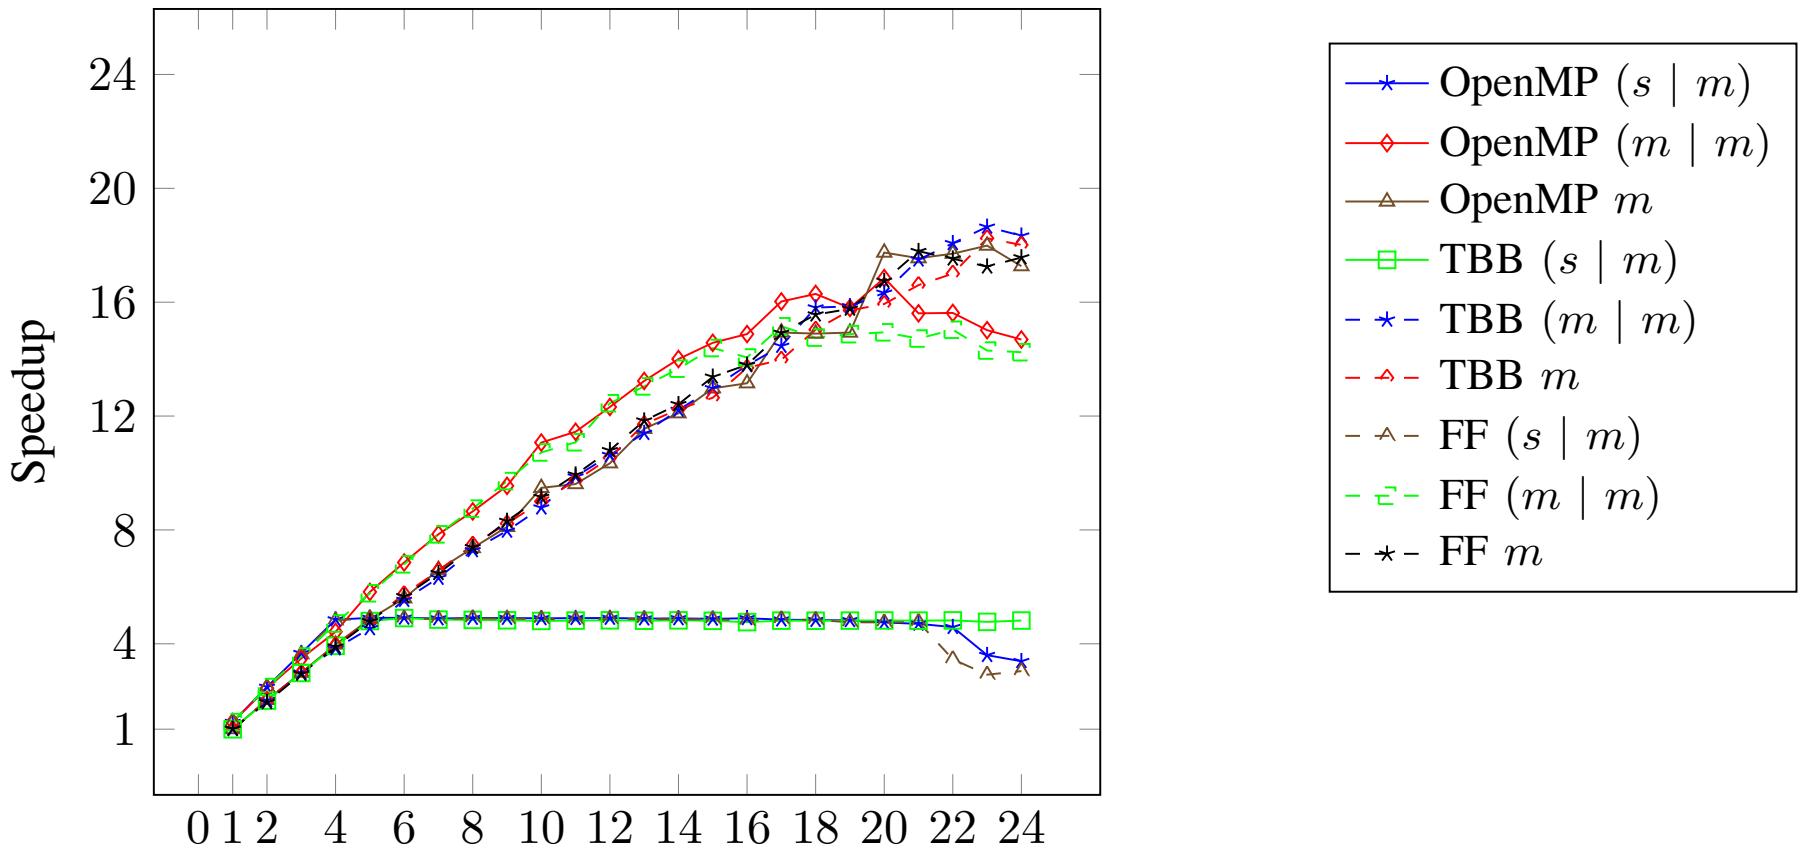
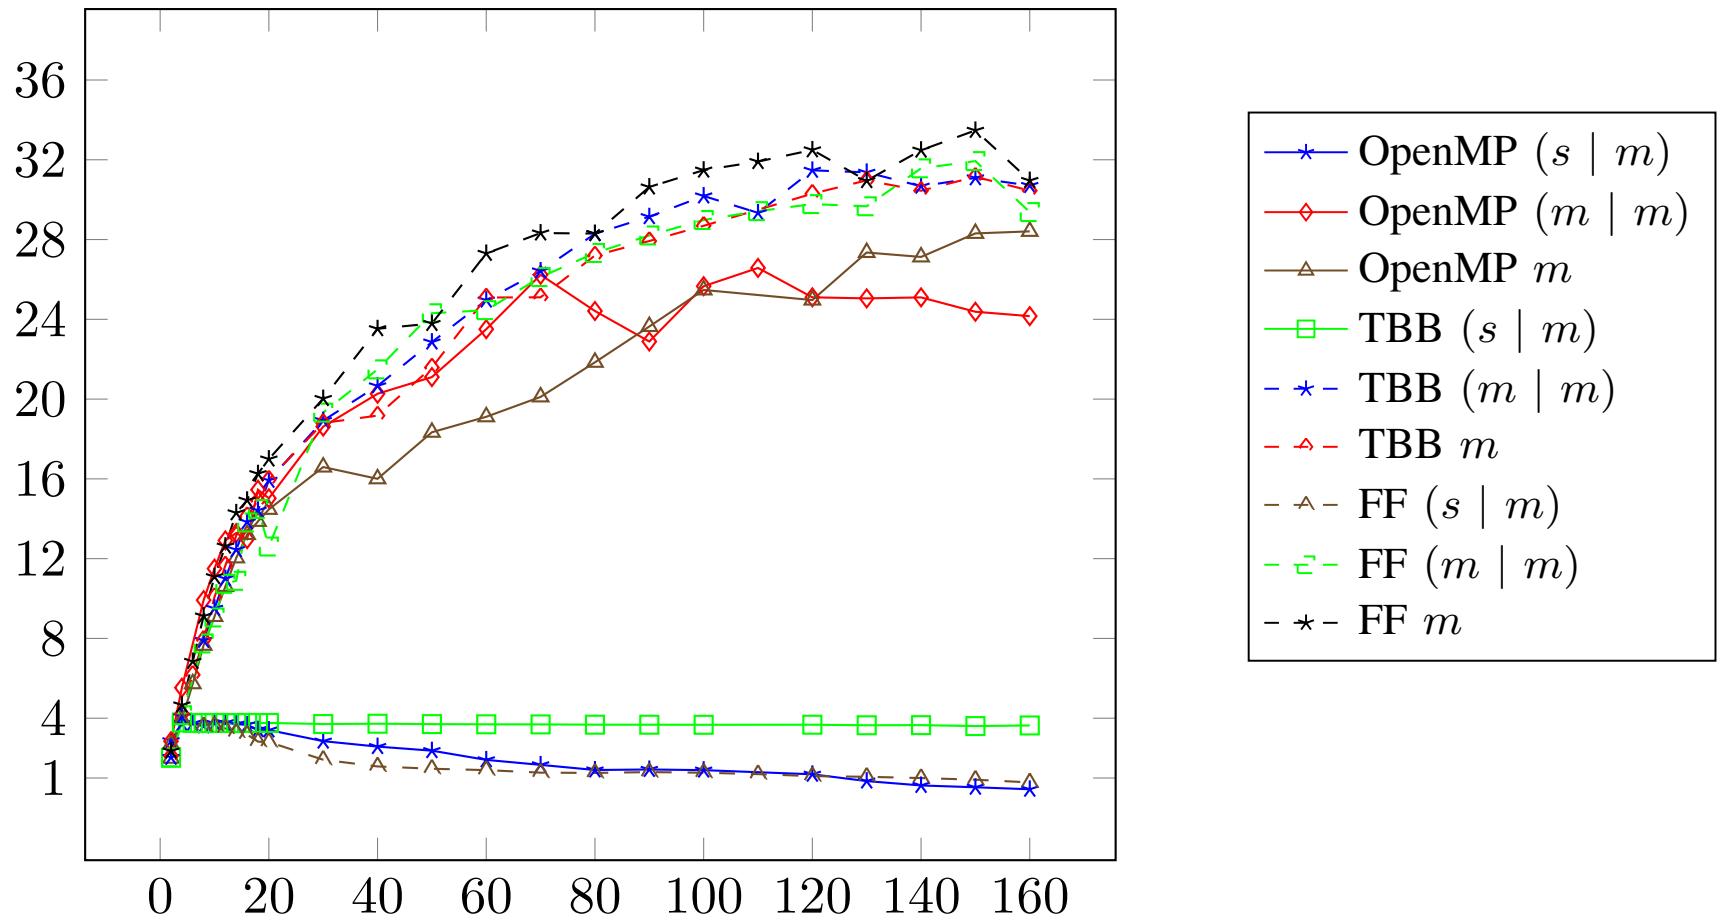


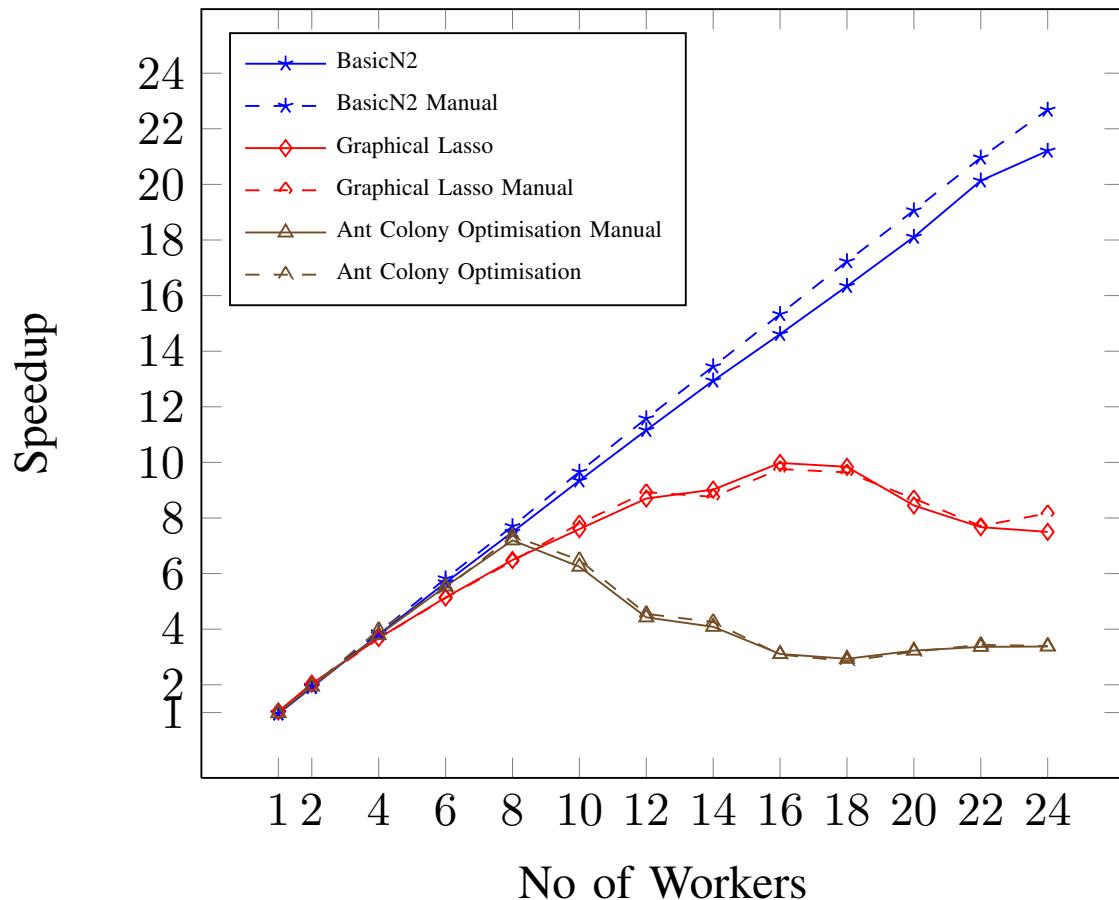
Image Convolution – 20 Cores!

Speedups for Image Convolution on *power8*



Comparable Performance

Speedups for Ant Colony, BasicN2 and Graphical Lasso

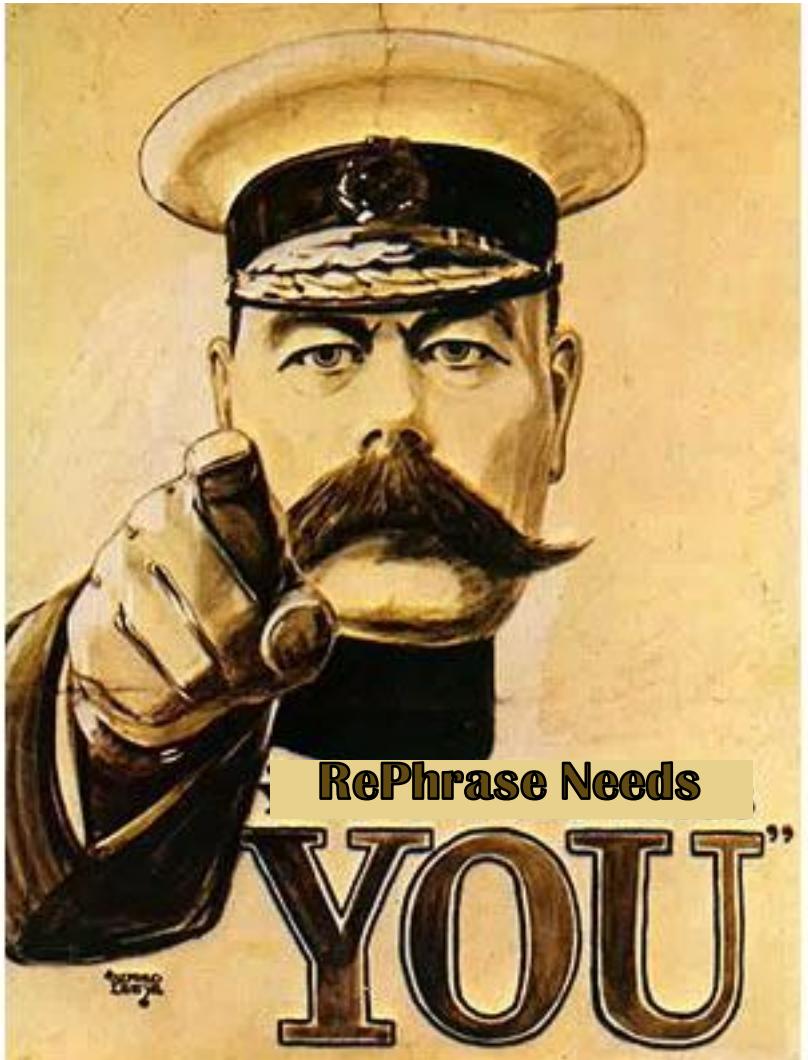


Conclusions

- Refactoring tool support:
 - Guides a programmer through steps to achieve parallelism
 - Warns the user if they are going wrong
 - Avoids common pitfalls
 - Helps with **understanding** and **intuition**
 - Reduces amount of boilerplate code
 - Allows programmer to concentrate on algorithm, rather than parallelism.

ParaPerformance/Rephrase Needs You!

- Please join our mailing list and help grow our user community
 - news items
 - access to free development software
 - chat to the developers
 - free developer workshops
 - bug tracking and fixing
 - Tools for both Erlang and C++
- Subscribe at
<https://mailman.cs.st-andrews.ac.uk/mailman/listinfo/rephrase-news>
- We're also looking for open source developers...





University
of
St Andrews

THANK YOU!

<http://rephrase-ict.eu>

@rephrase_eu